

Machine Learning


Graph Kernels

Manfred Jaeger


Aalborg University


Reminder


SVM learning after a non-linear transformation:



Problem: suitable feature spaces may need to be very high-dimensional. That makes computations with transformed vectors $\phi(\mathbf{x})$ very expensive.







➡ When we have a positive semi-definite kernel $K(\mathbf{x}, \mathbf{x}')$ we need not know the actual feature mapping ϕ

We obtain the following strategy to construct *Kernelized* SVM classifiers:

- ▶ **Given:** a classification problem
- ▶ **Define** or **Select** a similarity function

$$K(\mathbf{x}, \mathbf{z})$$

- ▶ **Verify** that $K(\mathbf{x}, \mathbf{z})$ is positive semi-definite
- ▶ **Learn** an SVM using the values $K(\mathbf{x}_i, \mathbf{x}_j)$ in place of dot products $\mathbf{x}_i \cdot \mathbf{x}_j$.

➡ The learning algorithm requires as input “only” the kernel matrix for the training datapoints $\mathbf{x}_1, \dots, \mathbf{x}_N$.

We obtain the following strategy to construct *Kernelized* SVM classifiers:

- ▶ **Given:** a classification problem
- ▶ **Define** or **Select** a similarity function

$$K(\mathbf{x}, \mathbf{z})$$

- ▶ **Verify** that $K(\mathbf{x}, \mathbf{z})$ is positive semi-definite
- ▶ **Learn** an SVM using the values $K(\mathbf{x}_i, \mathbf{x}_j)$ in place of dot products $\mathbf{x}_i \cdot \mathbf{x}_j$.

➡ The learning algorithm requires as input “only” the kernel matrix for the training datapoints $\mathbf{x}_1, \dots, \mathbf{x}_N$.

➡ This also opens the SVM technology to datatypes other than numeric data: just need to find kernels: positive semi-definite (similarity) functions on pairs of data objects.

Given: a set of graphs $\mathcal{G}$, e.g.

- ▶ $\mathcal{G}$: set of all finite undirected graphs (V, E)
- ▶ $\mathcal{G}$: set of all finite directed graphs (V, E, A_1, A_2) with node attributes A_1, A_2 .
- ▶ ...

Want: a feature function

$$\phi : G \mapsto \phi(G) \in \mathbb{R}^d \quad (G \in \mathcal{G})$$

(explicit construction), **or** a positive semi-definite kernel

$$K : (G_0, G_1) \mapsto K(G_0, G_1) \in \mathbb{R} \quad (G_0, G_1 \in \mathcal{G})$$

(implicit construction).

Convolution Kernels

Let $\mathcal{D}$ be a (structured) data type (e.g.: $\mathcal{D} = \mathcal{G}$ a type of graphs).

Decompositions

- ▶ $\mathcal{P}$: a space of “components” or “parts” of objects from $\mathcal{D}$
- ▶ $\mathcal{R}$: a relation between parts and objects with the interpretation: $(p, d) \in \mathcal{R}$ if p is a part of d .
- ▶ Define $\mathcal{R}^{-1}(d) := \{p | (p, d) \in \mathcal{R}\}$ (the parts of d)

Convolution Kernels

Let $K_{\mathcal{P}}$ be a Kernel on $\mathcal{P}$. Then

$$K_D(d, d') = \sum_{p \in \mathcal{R}^{-1}(d)} \sum_{p' \in \mathcal{R}^{-1}(d')} K_{\mathcal{P}}(p, p')$$

is a Kernel on $\mathcal{D}$.

Notes:

- ▶ more complex decompositions into different types of parts are usually considered
- ▶ the $\mathcal{R}$ relation is very abstract and need not correspond to an intuitive “part of” relationship.

[Haussler, David. Convolution kernels on discrete structures. Vol. 646. Technical report, Department of Computer Science, University of California at Santa Cruz, 1999.]

Dot product as convolution kernel

$$\mathcal{D} = \mathbb{R}^n, \mathcal{P} = \{1, \dots, n\} \times \mathbb{R}.$$

$$((i, z), \mathbf{x}) \in \mathcal{R} \Leftrightarrow \mathbf{x}[i] = z.$$

Parts kernel:

$$K_P((i, y), (j, z)) = \begin{cases} y \cdot z & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$$

Then

$$K_D(\mathbf{x}, \mathbf{y}) = \mathbf{x} \cdot \mathbf{y}.$$

Dot product as convolution kernel

$$\mathcal{D} = \mathbb{R}^n, \mathcal{P} = \{1, \dots, n\} \times \mathbb{R}.$$

$$((i, z), \mathbf{x}) \in \mathcal{R} \Leftrightarrow \mathbf{x}[i] = z.$$

Parts kernel:

$$K_P((i, y), (j, z)) = \begin{cases} y \cdot z & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$$

Then

$$K_D(\mathbf{x}, \mathbf{y}) = \mathbf{x} \cdot \mathbf{y}.$$

Vectors of different lengths

$$\mathcal{D} = \bigcup_{n \in \mathbb{N}} \mathbb{R}^n, \mathcal{P} = \mathbb{N} \times \mathbb{R}.$$

$$((i, z), \mathbf{x}) \in \mathcal{R} \Leftrightarrow \mathbf{x}[i] = z.$$

Parts kernel:

$$K_P((i, y), (j, z)) = y \cdot z.$$

Then for $\mathbf{x} \in \mathbb{R}^n, \mathbf{y} \in \mathbb{R}^m$:

$$K_D(\mathbf{x}, \mathbf{y}) = \sum_{i=1}^n \sum_{j=1}^m \mathbf{x}[i] \cdot \mathbf{y}[j].$$

➡ Convolution kernels enable comparison of objects of different sizes (without defining an “alignment” of parts).

We consider $\mathcal{D} = \Sigma^*$ for some alphabet Σ .

String Kernel 1

Define $\mathcal{P} = \Sigma^n$ (" n -grams") and for $\mathbf{u} \in \mathcal{P}$, $\mathbf{s} \in \mathcal{D}$:

$$(\mathbf{u}, \mathbf{s}) \in \mathcal{R} \Leftrightarrow \mathbf{u} \text{ appears as a substring in } \mathbf{s}.$$

Parts kernel:

$$K_{\mathcal{P}}(\mathbf{u}, \mathbf{v}) = \begin{cases} 0 & \text{if } \mathbf{u} \neq \mathbf{v} \\ 1 & \text{if } \mathbf{u} = \mathbf{v} \end{cases} \quad (\mathbf{u}, \mathbf{v} \in \mathcal{P})$$

Then:

$$K_{\mathcal{D}}(\mathbf{s}, \mathbf{t}) = |\{\mathbf{u} \in \Sigma^n \mid (\mathbf{u}, \mathbf{s}) \in \mathcal{R} \text{ and } (\mathbf{u}, \mathbf{t}) \in \mathcal{R}\}|$$

(number of common n -grams of $\mathbf{s}$ and $\mathbf{t}$).

String Kernel 2

Define $\mathcal{P} = \mathbb{N} \times \Sigma^n$ and for $(k, \mathbf{u}) \in \mathcal{P}$, $\mathbf{s} \in \mathcal{D}$:

$((k, \mathbf{u}), \mathbf{s}) \in \mathcal{R} \Leftrightarrow \mathbf{u}$ is the substring of length n starting at index k in $\mathbf{s}$.

Parts kernel:

$$K_P((k, \mathbf{u}), (l, \mathbf{v})) = \begin{cases} 0 & \text{if } \mathbf{u} \neq \mathbf{v} \\ 1 & \text{if } \mathbf{u} = \mathbf{v} \end{cases} \quad ((k, \mathbf{u}), (l, \mathbf{v}) \in \mathcal{P})$$

Then:

$$K_D(\mathbf{s}, \mathbf{t}) = \sum_{\mathbf{u} \in \Sigma^n} |\{k \in \mathbb{N} | ((k, \mathbf{u}), \mathbf{s}) \in \mathcal{R}\}| \cdot |\{k \in \mathbb{N} | ((k, \mathbf{u}), \mathbf{t}) \in \mathcal{R}\}|$$

➡ This is the n -spectrum kernel!

$\mathcal{D}$: finite graphs (V, E, L) with a node label attribute L .

$\mathcal{P}$: set of nodes in graphs from $\mathcal{D}$.

$$(v, (V, E, L)) \in \mathcal{R} \Leftrightarrow v \in V$$

Parts kernel: kernel on L , e.g. 0/1 for categorical L , product for numerical L .

Then:

$$K_D((V, E, L), (V', E', L')) = \sum_{v \in V} \sum_{v' \in V'} K_P(L(v), L(v')).$$

➡ Graph structures given by E, E' not taken into account!

Sub-graphs and Isomorphisms

Encoding E for a graph with vertices $V = \{1, \dots, n\}$:

Adjacency List

Directed:

list the targets of outgoing edges of a node:

1: 2
2: 6
3:
4: 2,6
5: 1,6
6: 3

Undirected:

list all neighbors of a node:

1: 2,3,5
2: 1,3
3: 1,2
4: 6
5: 1,6
6: 4,5

Adjacency Matrix


For each pair of vertices: indicate whether edge exists

Directed:


$$\begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 0 \end{pmatrix}$$

Undirected: symmetric matrix

$$\begin{pmatrix} 0 & 1 & 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 1 & 0 \end{pmatrix}$$




$$G_1 = (V_1, E_1)$$




$$G_2 = (V_2, E_2)$$

Are these two graphs identical?



$$G_1 = (V_1, E_1)$$



$$G_2 = (V_2, E_2)$$

Are these two graphs identical?

No, because

$$\begin{aligned} V_1 &= \{Coca\ Cola, Apple, Novo\ Nordisk, KLM\} \\ &\neq \\ V_2 &= \{Susan, Peter, Christopher, Mary\} \end{aligned}$$

Isomorphic Graphs

$(V_1, E_1), (V_2, E_2)$ are *isomorphic*, if there exists a bijective mapping

$$f : V_1 \rightarrow V_2$$

such that for all $v, v' \in V_1$:

$$(v, v') \in E_1 \Leftrightarrow (f(v), f(v')) \in E_2.$$

The function f then is called a *graph isomorphism*.

Isomorphic Graphs

$(V_1, E_1), (V_2, E_2)$ are *isomorphic*, if there exists a bijective mapping



$$f : V_1 \rightarrow V_2$$

such that for all $v, v' \in V_1$:

$$(v, v') \in E_1 \Leftrightarrow (f(v), f(v')) \in E_2.$$

The function f then is called a *graph isomorphism*.

Example



are isomorphic, as witnessed by:

$$f : \text{Coca Cola} \mapsto \text{Susan}, \text{Apple} \mapsto \text{Peter}, \text{Novo Nordisk} \mapsto \text{Christopher}, \text{KLM} \mapsto \text{Mary}$$



Question: is this the only isomorphism?

(V_1, E_1) is *subgraph-isomorphic* to (V_2, E_2) , if there exists an injective mapping

$$f : V_1 \rightarrow V_2$$

such that for all $v, v' \in V_1$:

$$(v, v') \in E_1 \Leftrightarrow (f(v), f(v')) \in E_2.$$




(V_1, E_1) is *subgraph-isomorphic* to (V_2, E_2) , if there exists an injective mapping

$$f : V_1 \rightarrow V_2$$

such that for all $v, v' \in V_1$:

$$(v, v') \in E_1 \Leftrightarrow (f(v), f(v')) \in E_2.$$




(V_1, E_1) is *subgraph-isomorphic* to (V_2, E_2) , if there exists an injective mapping

$$f : V_1 \rightarrow V_2$$

such that for all $v, v' \in V_1$:

$$(v, v') \in E_1 \Leftrightarrow (f(v), f(v')) \in E_2.$$



Problem

Given graphs (V_1, E_1) , (V_2, E_2) (as adjacency matrices, adjacency lists, ...), is (V_1, E_1) (sub-graph) isomorphic to (V_2, E_2) ?

Answer

Can check each of the

$$\frac{|V_2|!}{(|V_2| - |V_1|)!}$$

possible injective mappings $f : V_1 \rightarrow V_2$.

➡ This is infeasible for large graphs (exponential in $|V_1|$)!

Open Problem

Is there an efficient way to decide whether (V_1, E_1) is an isomorphic subgraph of (V_2, E_2) ?

➡ Probably not: this is an NP-hard problem!

Graphlet Kernel

For simplicity: consider undirected graphs without attributes.

Graphlet


Graphlet: a small graph (e.g. size 3,4,5). We consider collections $\mathcal{G}$ of graphlets, where usually

- ▶ all $g \in \mathcal{G}$ have the same size
- ▶ only consider connected g
- ▶ at most one representative of each isomorphism class included in $\mathcal{G}$

Size 3 graphlets:



Size 4 graphlets:




→ similar to n -grams for strings

[Shervashidze, Nino, et al. "Efficient graphlet kernels for large graph comparison." Artificial intelligence and statistics. PMLR, 2009.]

Graphlet features

Graphlet g defines feature ϕ_g of (bigger) graph G : number of occurrences of g as a subgraph of G .


Version 1: count the number of isomorphisms that map g to an induced subgraph of G :



Graphlet features

Graphlet g defines feature ϕ_g of (bigger) graph G : number of occurrences of g as a subgraph of G .


Version 1: count the number of isomorphisms that map g to an induced subgraph of G :



Graphlet features

Graphlet g defines feature ϕ_g of (bigger) graph G : number of occurrences of g as a subgraph of G .


Version 1: count the number of isomorphisms that map g to an induced subgraph of G :



Graphlet features

Graphlet g defines feature ϕ_g of (bigger) graph G : number of occurrences of g as a subgraph of G .

Version 1: count the number of isomorphisms that map g to an induced subgraph of G :




6 different isomorphisms embed g_1 in G : $\phi_{g_1}^{(1)}(G) = 6$

Graphlet features

Graphlet g defines feature ϕ_g of (bigger) graph G : number of occurrences of g as a subgraph of G .

Version 1: count the number of isomorphisms that map g to an induced subgraph of G :




$4! = 24$ different isomorphisms embed g_2 in G : $\phi_{g_2}^{(1)}(G) = 24$

Graphlet features

Graphlet g defines feature ϕ_g of (bigger) graph G : number of occurrences of g as a subgraph of G .

Version 1: count the number of isomorphisms that map g to an induced subgraph of G :



$4! = 24$ different isomorphisms embed g_2 in G : $\phi_{g_2}^{(1)}(G) = 24$

Version 2: only count how many induced subgraphs are isomorphic to $g \in \mathcal{G}$:

$$\phi_{g_1}^{(2)}(G) = 3, \phi_{g_2}^{(2)}(G) = 1.$$

Raw count feature vectors $\phi^{(1)}$ or $\phi^{(2)}$ have larger entries for larger graphs.

➡ according to the corresponding kernels $K^{(1)}, K^{(2)}$, larger graphs will be “more similar” to other graphs than smaller graphs.



Solution 1: normalize the feature vectors ϕ to probability distributions.

Solution 2: apply the generic kernel normalization (normalize feature vectors to length 1):



$$K(G, G') \rightarrow \frac{K(G, G')}{\sqrt{K(G, G) \cdot K(G', G')}}.$$

Random Walk Kernels

Node-labeled graphs (there could also be edge labels):



Node-labeled graphs (there could also be edge labels):



Random walk:

- Pick a random start vertex (e.g. uniform)
- make random transitions (e.g. uniformly to any neighbor)
- terminate random walk at any step with a fixed termination probability p_t

Label sequence	A Probability
C	$\frac{3}{5}p_t$
O	$\frac{1}{5}p_t$
CO	$\frac{1}{5}\frac{1}{1}p_t + \frac{1}{5}\frac{1}{2}p_t + \frac{1}{5}\frac{1}{2}p_t$
...	...

Label sequence	B Probability
C	$\frac{2}{4}p_t$
O	$\frac{1}{4}p_t$
CO	$\frac{1}{4}\frac{1}{2}p_t$
...	...

Label sequence kernel

Let Σ be the alphabet of node labels. The random walk model defines for each graph G a probability distribution p_G over Σ^* . This is the feature vector for the *label sequence kernel*:

$$K(G, G') = \sum_{\mathbf{s} \in \Sigma^*} p_G(\mathbf{s}) p_{G'}(\mathbf{s})$$

“Efficient” computation of this infinite sum possible in $O(n^4)$ by dynamic programming (n : max. number of nodes in G or G')!

Adding string kernel

Taking similarity of distinct label sequences into account:

$$K(G, G') = \sum_{\mathbf{s} \in \Sigma^*} \sum_{\mathbf{t} \in \Sigma^*} K(\mathbf{s}, \mathbf{t}) p_G(\mathbf{s}) p_{G'}(\mathbf{t})$$

where $K(\mathbf{s}, \mathbf{t})$ could be any string kernel.


Including edge labels

Random walk kernels easily incorporate edge labels: random walks then are sequences of alternating node and edge labels.


[Tsuda, Saigo: Graph Classification. In: Aggarwal et al.: Managing and Mining Graph Data. Springer 2010]

Weisfeiler-Lehman Kernel

W-L node code refinement process for (undirected) graph (V, E)




W-L node code refinement process for (undirected) graph (V, E)




Initialize node codes c^0 : node attribute vector if exists;
otherwise a constant code

W-L node code refinement process for (undirected) graph (V, E)



Initialize node codes c^0 : node attribute vector if exists;
otherwise a constant code

W-L node code refinement process for (undirected) graph (V, E)




Initialize node codes c^0 : node attribute vector if exists; otherwise a constant code

In iteration k : compute code

$$c^k(v) = \text{hash}((c^{k-1}(v), \{c^{k-1}(w) : (v, w) \in E\}))$$

where hash is a function that takes as input a pair consisting of a code and a multiset of codes, and returns a unique new code.

W-L node code refinement process for (undirected) graph (V, E)




Initialize node codes c^0 : node attribute vector if exists; otherwise a constant code

In iteration k : compute code

$$c^k(v) = \text{hash}((c^{k-1}(v), \{c^{k-1}(w) : (v, w) \in E\}))$$

where hash is a function that takes as input a pair consisting of a code and a multiset of codes, and returns a unique new code.

W-L node code refinement process for (undirected) graph (V, E)




Initialize node codes c^0 : node attribute vector if exists; otherwise a constant code

In iteration k : compute code

$$c^k(v) = \text{hash}((c^{k-1}(v), \{c^{k-1}(w) : (v, w) \in E\}))$$

where hash is a function that takes as input a pair consisting of a code and a multiset of codes, and returns a unique new code.

W-L node code refinement process for (undirected) graph (V, E)




Initialize node codes c^0 : node attribute vector if exists; otherwise a constant code

In iteration k : compute code

$$c^k(v) = \text{hash}((c^{k-1}(v), \{c^{k-1}(w) : (v, w) \in E\}))$$

where hash is a function that takes as input a pair consisting of a code and a multiset of codes, and returns a unique new code.

W-L node code refinement process for (undirected) graph (V, E)



Initialize node codes c^0 : node attribute vector if exists; otherwise a constant code

In iteration k : compute code

$$c^k(v) = \text{hash}((c^{k-1}(v), \{c^{k-1}(w) : (v, w) \in E\}))$$

where hash is a function that takes as input a pair consisting of a code and a multiset of codes, and returns a unique new code.

- ▶ At iteration k the codes define a partition P^k of the nodes into sets of nodes with equal codes
- ▶ P^k is a refinement of P^{k-1}
- ▶ If $P^k = P^{k-1}$, then $P^{k+l} = P^{k-1}$ for all $l \geq 0$
- ▶ After at most $|V|$ iterations the partition becomes stable

Let $G = (V, E)$, $G' = (V', E')$ be isomorphic graphs with isomorphism $I : V \rightarrow V'$. Then for all $v \in V, k \geq 0$:

$$c^k(I(v)) = c^k(v)$$

(proof by induction on k). In particular:

$$\{c^k(v) : v \in V\} = \{c^k(v') : v' \in V'\}$$

Let $G = (V, E)$, $G' = (V', E')$ be isomorphic graphs with isomorphism $I : V \rightarrow V'$. Then for all $v \in V, k \geq 0$:

$$c^k(I(v)) = c^k(v)$$

(proof by induction on k). In particular:

$$\{c^k(v) : v \in V\} = \{c^k(v') : v' \in V'\}$$

WL-test for isomorphism: given graphs G, G' with $n = |V| = |V'|$:

- ▶ run WL code updates for n iterations;
- ▶ if $\{c^n(v) : v \in V\} \neq \{c^n(v') : v' \in V'\}$: return “not isomorphic”; otherwise: return “don’t know”.

Let $G = (V, E)$, $G' = (V', E')$ be isomorphic graphs with isomorphism $I : V \rightarrow V'$. Then for all $v \in V, k \geq 0$:

$$c^k(I(v)) = c^k(v)$$




(proof by induction on k). In particular:

$$\{c^k(v) : v \in V\} = \{c^k(v') : v' \in V'\}$$

WL-test for isomorphism: given graphs G, G' with $n = |V| = |V'|$:

- ▶ run WL code updates for n iterations;
- ▶ if $\{c^n(v) : v \in V\} \neq \{c^n(v') : v' \in V'\}$: return “not isomorphic”; otherwise: return “don’t know”.

Exercise: which of the following graphs will be identified as non-isomorphic by the WL-test?



Let $\mathcal{C}^k$ be the code space for codes generated in the k 'th iteration (for all possible graphs).

Each $\kappa \in \mathcal{C}^k$ defines a feature for graphs $G = (V, E)$:

$$\phi^\kappa(G) = |\{v \in V : c^k(v) = \kappa\}|$$

Fix $h \geq 0$. The **WL-kernel** of depth h is defined by the features

$$\{\phi^\kappa : \kappa \in \mathcal{C}^k, k = 0, \dots, h\}$$

→ infinite dimensional feature vectors ($\mathcal{C}^k$ infinite for $k \geq 1$), but for any graph G only finitely many $\phi^\kappa(G)$ are nonzero.

→ more precisely, this is the WL **subtree kernel**. There are some other types of WL-kernels.

Kernels for node classification


If we want a kernel for nodes, rather than full graphs:

Node centered random walks

Feature vector for node $v \in G$ given by label sequence probabilities for random walks starting at v

Ego graphs

For any graph kernel $k(G, G')$: can apply kernel to the *ego graphs* of nodes v (subgraph consisting of v and its immediate neighbors).



⚡ beware: the node of interest has no special role in the application of $k(G, G')$ to the ego graphs